

GTU Department of Computer Engineering

CSE 312 – Spring 2024

HW 1 Report

Bariş Batuhan Bolat

210104004029

HW 1 Part A

- **syscalls.cpp**

- **getPid()** : Gets the process ID
- **waitpid()** : Waits for a child process with the given ID
- **sys_exit()** : Exits the current process
- **fork()** : Creates a child process (without returning the child's PID)
- **fork(int *pid)** : Creates a child process and returns the child's PID in the provided pointer
- **exec(void entrypoint())** : Executes and replaces current task.
- **addTask(void entrypoint())** : Adds a new task to the system
- **printf** : Prints the desired char sequence to screen. Implementation is in **kernel.cpp**.

```
// Gets the process ID
int myos::getPid()
{
    int pId = 1;
    asm("int $0x80" : "=c"(pId) : "a"(SYSCALLS::GETPID));
    return pId;
}

// Waits for a child process with the given ID
void myos::waitpid(common::uint8_t wPid)
{
    asm("int $0x80" : : "a"(SYSCALLS::WAITPID), "b"(wPid));
}

// Exits the current process
void myos::sys_exit()
{
    asm("int $0x80" : : "a"(SYSCALLS::EXIT));
}

// Creates a child process (without returning the child's PID)
void myos::fork()
{
    asm("int $0x80" : : "a"(SYSCALLS::FORK));
}

// Creates a child process and returns the child's PID in the provided pointer
void myos::fork(int *pid)
{
    asm("int $0x80" : "=c"(*pid) : "a"(SYSCALLS::FORK));
}

int myos::exec(void entrypoint())
{
    int result;
    asm("int $0x80" : "=c"(result) : "a"(SYSCALLS::EXEC), "b"((uint32_t)entrypoint));
    return result;
}

// Adds a new task to the system
int myos::addTask(void entrypoint())
{
    int result;
    asm("int $0x80" : "=c"(result) : "a"(SYSCALLS::ADDTASK), "b"((uint32_t)entrypoint));
    return result;
}
```

```
enum SYSCALLS
{
    EXIT, GETPID, WAITPID, FORK, EXEC, PRINTF, ADDTASK
};
```

- **Syscalls HandleInterrupt**

Acts as a bridge between user-space system call requests and the kernel's system call implementation a.k.a interrupts.

```
uint32_t SyscallHandler::HandleInterrupt(uint32_t esp)
{
    // Get a pointer to the CPU state structure from the stack pointer
    CPUState *cpu = (CPUState *)esp;
    // Switch statement to handle different system calls based on the value in cpu->eax
    switch (cpu->eax)
    {
        case SYSCALLS::EXEC:
            esp = InterruptHandler::sys_exec(cpu->ebx);
            break;
        case SYSCALLS::FORK:
            cpu->ecx = InterruptHandler::sys_fork(cpu);
            return InterruptHandler::HandleInterrupt(esp);
            break;
        case SYSCALLS::PRINTF:
            printf((char *)cpu->ebx);
            break;
        case SYSCALLS::EXIT:
            if (InterruptHandler::sys_exit())
            {
                return InterruptHandler::HandleInterrupt(esp);
            }
            break;
        case SYSCALLS::WAITPID:
            if (InterruptHandler::sys_waitpid(esp))
            {
                return InterruptHandler::HandleInterrupt(esp);
            }
            break;
        case SYSCALLS::GETPID:
            cpu->ecx = InterruptHandler::sys_getpid();
            break;
        case SYSCALLS::ADDTASK:
            cpu->ecx = InterruptHandler::sys_addTask(cpu->ebx);
            break;
        default:
            break;
    }
    return esp;
}
```

- **interrupts.cpp(Changes)**

- **sys_getpid()** : Gets the process ID of the current task by calling task managers **GetPid**.
- **sys_exec(common::uint32_t entrypoint)** : Executes a new task by calling task managers **ExecTask**.
- **sys_addTask(common::uint32_t entrypoint)** : Adds a new task by calling task managers **AddTask**.
- **sys_fork(CPUState *cpustate)** : Forks the current task by calling task managers **ForkTask**.
- **sys_exit()** : Exits the current task by calling task managers **ExitCurrentTask**.
- **sys_waitpid(common::uint32_t pid)** : Wait for a child task to exit by calling task managers **WaitTask**.

```

common::uint32_t InterruptHandler::sys_getpid()
{
    return interruptManager->taskManager->GetPid();
}

common::uint32_t InterruptHandler::sys_exec(common::uint32_t entrypoint)
{
    return interruptManager->taskManager->ExecTask((void (*)())entrypoint);
}

common::uint32_t InterruptHandler::sys_addTask(common::uint32_t entrypoint)
{
    return interruptManager->taskManager->AddTask((void (*)())entrypoint);
}

common::uint32_t InterruptHandler::sys_fork(CPUState *cpustate)
{
    return interruptManager->taskManager->ForkTask(cpustate);
}

bool InterruptHandler::sys_exit()
{
    return interruptManager->taskManager->ExitCurrentTask();
}

bool InterruptHandler::sys_waitpid(common::uint32_t pid)
{
    return interruptManager->taskManager->WaitTask(pid);
}

```

- Interrupt handlers **HandleInterrupt** function calls task manager scheduler function **Schedule**.

```

uint32_t InterruptHandler::HandleInterrupt(uint32_t esp)
{
    return (uint32_t)interruptManager->taskManager->Schedule((CPUState *)esp);
}

```

- **multitasking.cpp**

- **ForkTask(CPUState *cpustate) :** ForkTask creates a new process.
 - It checks for available slots and copies the parent task's state.
 - Increases pId by using static pIdCounter variable.
 - It copies current tasks stack to new created one.
 - On success, it returns the new task's process ID (pId).

```
common::uint32_t TaskManager::ForkTask(CPUState *cpustate)
{
    if (taskCount >= 256)
    {
        printf("There are 256 processes are active. No more can be created.\n");
        return 0;
    }
    tasks[taskCount].taskState = READY;
    tasks[taskCount].pId = tasks[currentTask].pId;
    tasks[taskCount].pId = (Task::pIdCounter)++;

    // Copy the parent task's stack to the new task's stack
    for (int i = 0; i < sizeof(tasks[currentTask].stack); i++)
    {
        tasks[taskCount].stack[i] = tasks[currentTask].stack[i];
    }

    // Calculate the offset of the CPU state within the parent task's stack
    common::uint32_t currentTaskOffset = (((common::uint32_t)cpustate - (common::uint32_t)tasks[currentTask].stack));

    // Set the new task's CPU state pointer based on the offset and its own stack
    tasks[taskCount].cpustate = (CPUState *)(((common::uint32_t)tasks[taskCount].stack) + currentTaskOffset);

    // Initialize specific CPU register (ECX) to 0
    tasks[taskCount].cpustate->ecx = 0;
    taskCount++;

    // Set the new task's time slice (timeRemaining)
    tasks[taskCount].timeRemaining = timeQuantum;
    return tasks[taskCount - 1].pId;
}
```

- **AddTask(Task *task) :** Adds a new task with the provided task object.
 - It checks for available slots and copies the state, cpu registers, program counter and specific flags from the provided task.
 - Sets the time of the new task with timeQuantum.
 - On success, it returns true.

```
bool TaskManager::AddTask(Task *task)
{
    if (taskCount >= 256)
    {
        return false;
    }
    tasks[taskCount].taskState = READY;
    tasks[taskCount].pId = task->pId;
    // Set the new task's CPU state pointer to the top of its stack with an offset for CPUState size
    tasks[taskCount].cpustate = (CPUState *) (tasks[taskCount].stack + 4096 - sizeof(CPUState));

    // Copy specific CPU registers (EAX, EBX, ECX, EDX) from the provided task
    tasks[taskCount].cpustate->eax = task->cpustate->eax;
    tasks[taskCount].cpustate->ebx = task->cpustate->ebx;
    tasks[taskCount].cpustate->ecx = task->cpustate->ecx;
    tasks[taskCount].cpustate->edx = task->cpustate->edx;

    // Copy additional CPU registers (ESI, EDI, EBP) from the provided task
    tasks[taskCount].cpustate->esi = task->cpustate->esi;
    tasks[taskCount].cpustate->edi = task->cpustate->edi;
    tasks[taskCount].cpustate->ebp = task->cpustate->ebp;

    // Copy program counter (EIP) and code segment (CS) from the provided task
    tasks[taskCount].cpustate->eip = task->cpustate->eip;
    tasks[taskCount].cpustate->cs = task->cpustate->cs;

    // Copy status flags (EFLAGS) from the provided task
    tasks[taskCount].cpustate->eflags = task->cpustate->eflags;
    // Set the new task's time slice (timeRemaining)
    tasks[taskCount].timeRemaining = timeQuantum;
    taskCount++;
    return true;
}
```

- **AddTask(void entrypoint())** : Adds a new task with the provided task object.
 - It checks for available slots and copies the state, cpu registers, program counter and specific flags from the provided entry point function.
 - Sets the time of the new task with timeQuantum.
 - On success, it returns true.

```
common::uint32_t TaskManager::AddTask(void entrypoint())
{
    tasks[taskCount].taskState = READY;
    tasks[taskCount].pId = (Task::pIdCounter)++;
    // Set the new task's CPU state pointer to the top of its stack with an offset for CPUState size
    tasks[taskCount].cpustate = (CPUState *) (tasks[taskCount].stack + 4096 - sizeof(CPUState));

    // Initialize specific CPU registers (EAX, EBX, ECX, EDX) to 0
    tasks[taskCount].cpustate->eax = 0;
    tasks[taskCount].cpustate->ebx = 0;
    tasks[taskCount].cpustate->ecx = 0;
    tasks[taskCount].cpustate->edx = 0;

    // Initialize additional CPU registers (ESI, EDI, EBP) to 0
    tasks[taskCount].cpustate->esi = 0;
    tasks[taskCount].cpustate->edi = 0;
    tasks[taskCount].cpustate->ebp = 0;

    // Set the program counter (EIP) to the entry point function address
    tasks[taskCount].cpustate->eip = (uint32_t)entrypoint;
    // Set the code segment selector from the global descriptor table (gdt)
    tasks[taskCount].cpustate->cs = gdt->CodeSegmentSelector();
    // Set specific status flags (EFLAGS)
    tasks[taskCount].cpustate->eflags = 0x202;
    taskCount++;
    tasks[taskCount].timeRemaining = timeQuantum;
    return tasks[taskCount - 1].pId;
}
```

- **WaitTask(common::uint32_t esp)** : Suspends the current task until another task with the specified process ID (pid) finishes.
 - It prevents a task from waiting for itself and checks for invalid targets or finished tasks.
 - Set the current task's CPU state pointer
 - Set the current task's state to WAITING.
 - Set the process ID to wait.
 - On success, it returns true.

```
bool TaskManager::WaitTask(common::uint32_t esp)
{
    CPUState *cpustate = (CPUState *)esp;
    // Extract the process ID to wait for from a CPU register (EBX)
    common::uint32_t pid = cpustate->ebx;
    if (tasks[currentTask].pId == pid || pid == 0) // prevention self waiting
        return false;
    int index = getIndex(pid);
    // Check if the target task exists and isn't already finished
    if (index == -1)
        return false;
    if (taskCount <= index || tasks[index].taskState == FINISHED)
        return false;
    tasks[currentTask].cpustate = cpustate;
    tasks[currentTask].waitPid = pid;
    tasks[currentTask].taskState = WAITING;
    return true;
}
```

- **ExecTask(void entrypoint())** : Replaces the current task's state with the provided entry point function(Task).
 - Set the current task's CPU state pointer to the top of its stack with an offset for CPUState size.
 - Initialize specific and additional CPU registers, program counter, code segment selector and flags.
 - Returns a pointer to the CPU state for context switching.

```
common::uint32_t TaskManager::ExecTask(void entrypoint())
{
    tasks[currentTask].taskState = READY;
    // Set the current task's CPU state pointer to the top of its stack with an offset for CPUState size
    tasks[currentTask].cpustate = (CPUState *) (tasks[currentTask].stack + 4096 - sizeof(CPUState));

    // Initialize specific CPU registers (EAX, EBX, ECX, EDX) to 0
    tasks[currentTask].cpustate->eax = 0;
    tasks[currentTask].cpustate->ebx = 0;
    tasks[currentTask].cpustate->ecx = tasks[currentTask].pId;
    tasks[currentTask].cpustate->edx = 0;

    // Initialize additional CPU registers (ESI, EDI, EBP) to 0
    tasks[currentTask].cpustate->esi = 0;
    tasks[currentTask].cpustate->edi = 0;
    tasks[currentTask].cpustate->ebp = 0;

    // Set the program counter (EIP) to the entry point function address
    tasks[currentTask].cpustate->eip = (uint32_t)entrypoint;
    // Set the code segment selector from the global descriptor table (gdt)
    tasks[currentTask].cpustate->cs = gdt->CodeSegmentSelector();
    // Set specific status flags (EFLAGS)
    tasks[currentTask].cpustate->eflags = 0x202;
    return (uint32_t)tasks[currentTask].cpustate;
}
```

- **ExitCurrentTask()** : Marks the current task as finished.
 - Prints process table when task is ended.

```
bool TaskManager::ExitCurrentTask()
{
    tasks[currentTask].taskState = FINISHED;
    PrintProcessTable();
    return true;
}
```

- **GetPID()**: Returns the process ID of the current task.

```
common::uint32_t TaskManager::GetPID()
{
    return tasks[currentTask].pId;
}
```

- **getIndex(common::uint32_t pid)** : Finds the index of a task based on its process ID (pId).

```
int TaskManager::getIndex(common::uint32_t pid)
{
    int index = -1;
    for (int i = 0; i < taskCount; i++)
    {
        if (tasks[i].pId == pid)
        {
            index = i;
            break;
        }
    }
    return index;
}
```

- **PrintProcessTable** : Prints process table with their pid, parent pid and their states

```
void TaskManager::PrintProcessTable()
{
    printf("\n^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^\n");
    printf("PID\tPPID\tSTATE\n");
    for (int i = 1; i < taskCount; i++)
    {
        printNum(tasks[i].pId);
        printf("\t ");
        printNum(tasks[i].pPid);
        printf("\t ");
        if (tasks[i].taskState == TaskState::READY)
        {
            if (i == currentTask)
                printf("RUNNING");
            else
                printf("READY");
        }
        else if (tasks[i].taskState == TaskState::WAITING)
        {
            printf("WAITING");
        }
        else if (tasks[i].taskState == TaskState::FINISHED)
        {
            printf("FINISHED");
        }
        printf("\n");
    }
}
```


- **Scheduler** : Round-robin scheduling algorithm for tasks. It selects the next task to run based on the following logic:
 - **Check for tasks:** It first verifies if there are any tasks available for scheduling (checks task count). If not, it simply returns the current CPU state.
 - **Save current task state:** If tasks exist, it saves the CPU state of the currently running task (if any).
 - **Initialize next task search:** It starts searching for the next ready task by setting an index to the task after the current one (wrapping around if needed).
 - **Loop for ready task:** It iterates through tasks until a task in the "**READY**" state is found.
 - **Waiting tasks:** During this loop, it checks for tasks waiting on another task to finish (state: WAITING). If a waiting task is encountered, it searches for the task it's waiting on (using waitPid).
 - **Waited-on task finished:** If the waited-on task is finished, the waiting task's state is updated to "**READY**" and it is removed from the waiting list.
 - **Wrap around:** The search for a ready task is circular, meaning it wraps back to the beginning of the task list if no ready tasks are found after iterating through all tasks.
 - **Break on loop back:** The loop breaks if it reaches the current task index again, indicating no ready tasks were found.
 - **Handle current task time:** If the current task was in a "**READY**" state, it reduces its remaining time quantum.
 - **Time quantum expired:** If the time quantum reaches zero, the current task's time is reset, and the scheduler moves to the next task.
 - **Move to next task (default):** If the current task wasn't ready or its time quantum didn't expire, the scheduler simply moves to the next task.
 - **Final loop for ready task (optional):** In case the initial loop for a ready task didn't find one due to waiting tasks becoming ready later, it does another loop with the same logic as step 4.
 - **Update current task and return:** Finally, it updates the current task index to the one found ready and returns the CPU state of the newly selected task for context switching.

```

CPUState *TaskManager::Schedule(CPUState *cpustate)
{
    // Check if there are no tasks to schedule
    if (taskCount <= 0)
    {
        return cpustate;
    }

    // Save the CPU state of the current task
    if (currentTask >= 0)
    {
        tasks[currentTask].cpustate = cpustate;
    }

    // Initialize the next task index to check
    int findTask = (currentTask + 1) % taskCount;
    // Loop until a ready task is found
    while (tasks[findTask].taskState != READY)
    {
        // Check for tasks waiting on another task to finish
        if (tasks[findTask].taskState == WAITING && tasks[findTask].waitPid > 0)
        {
            int waitTaskIndex = getIndex(tasks[findTask].waitPid);

            // If the waited-on task is done, update waiting task and continue
            if (waitTaskIndex > -1 && tasks[waitTaskIndex].taskState != WAITING)
            {
                if (tasks[waitTaskIndex].taskState == FINISHED)
                {
                    tasks[findTask].waitPid = 0;
                    tasks[findTask].taskState = READY;
                }
            }
        }

        // Move to the next task (circular fashion)
        findTask = (findTask + 1) % taskCount;
        // Break if we looped back to the current task (no ready tasks)
        if (findTask == currentTask)
        {
            break;
        }
    }
}

```

```

if (tasks[currentTask].taskState == READY)
{
    // If the task finishes its time quantum, reset and move to next task
    tasks[currentTask].timeRemaining -= timeQuantum;
    if (tasks[currentTask].timeRemaining <= 0)
    {
        tasks[currentTask].timeRemaining = timeQuantum;
        findTask = (currentTask + 1) % taskCount;
    }
}
// If current task wasn't ready, just move to the next one
else
{
    findTask = (currentTask + 1) % taskCount;
}

// Loop again until a ready task is found (in case previous loop didn't)
while (tasks[findTask].taskState != READY)
{
    if (tasks[findTask].taskState == WAITING && tasks[findTask].waitPid > 0)
    {
        int waitTaskIndex = getIndex(tasks[findTask].waitPid);
        if (waitTaskIndex > -1 && tasks[waitTaskIndex].taskState != WAITING)
        {
            if (tasks[waitTaskIndex].taskState == FINISHED)
            {
                tasks[findTask].waitPid = 0;
                tasks[findTask].taskState = READY;
            }
        }
    }
    findTask = (findTask + 1) % taskCount;
    if (findTask == currentTask)
    {
        break;
    }
}

// Update the current task index and return its CPU state
currentTask = findTask;
return tasks[currentTask].cpustate;
}

```

- **kernel.cpp**

- **printNum** : Converts the provided integer to char sequence and prints.

```
void printNum(int num)
{
    char numberStr[10];
    itoa(num, numberStr, 10);
    printf(numberStr);
}
```

- **printf** : Prints the provided char sequence to the screen

```
void printf(char *str)
{
    static uint16_t *VideoMemory = (uint16_t *)0xb8000;

    static uint8_t x = 0, y = 0;

    for (int i = 0; str[i] != '\0'; ++i)
    {
        switch (str[i])
        {
            case '\n':
                x = 0;
                y++;
                break;
            default:
                VideoMemory[80 * y + x] = (VideoMemory[80 * y + x] & 0xFF00) | str[i];
                x++;
                break;
        }

        if (x >= 80)
        {
            x = 0;
            y++;
        }

        if (y >= 25)
        {
            for (y = 0; y < 25; y++)
                for (x = 0; x < 80; x++)
                    VideoMemory[80 * y + x] = (VideoMemory[80 * y + x] & 0xFF00) | ' ';
            x = 0;
            y = 0;
        }
    }
}
```

- **printArray** : Prints the array with n integer elements to the screen.

```
void printArray(int arr[], int n)
{
    printf("Array : {");
    for (int i = 0; i < n; i++)
    {
        printNum(arr[i]);
        if (i + 1 != n)
        {
            printf(",");
        }
    }
    printf("} ");
}
```

- **forExample**

- The code utilizes a loop to call `fork` six times, each time storing the returned value in a corresponding `pid` variable.
- The return value of `fork` determines the execution path:
 - **Parent Process (return value > 0):**
 - The parent process continues execution after the `fork` call.
 - It uses `waitpid` to wait for the child process (identified by `pid`) to finish before proceeding.
 - **Child Process (return value == 0):**
 - The child process executes the code within the `if` block.
 - It prints a message indicating its purpose using `printf`.
 - It calls custom functions like `TaskCollatz1`, `TaskCollatz2`, `TaskCollatz3`, and `long_running_program`. These functions likely perform specific tasks.
 - Finally, the child process exits using `sys_exit`.

```
void forkExample()
{
    int pid1 = 0;
    int pid2 = 0;
    int pid3 = 0;
    int pid4 = 0;
    int pid5 = 0;
    int pid6 = 0;

    int parentPid = getpid();
    printf("Task Pid:");
    printNum(parentPid);
    printf("\n");

    fork(&pid1);
    if (pid1 == 0)
    {
        printf("Collatz_1 : ");
        TaskCollatz1();
        sys_exit();
    }
    waitpid(pid1);

    fork(&pid2);
    if (pid2 == 0)
    {
        printf("Collatz_2 : ");
        TaskCollatz2();
        sys_exit();
    }
    waitpid(pid2);

    fork(&pid3);
    if (pid3 == 0)
    {
        printf("Collatz_3 : ");
        TaskCollatz3();
        sys_exit();
    }
    waitpid(pid3);

    fork(&pid4);
    if (pid4 == 0)
    {
        printf("LongRunning_10000 : ");
        printNum(long_running_program(14000));
        printf("\n");
        sys_exit();
    }
    waitpid(pid4);

    fork(&pid5);
    if (pid5 == 0)
    {
        printf("LongRunning_20000 : ");
        printNum(long_running_program(18000));
        printf("\n");
        sys_exit();
    }
    waitpid(pid5);

    fork(&pid6);
    if (pid6 == 0)
    {
        printf("LongRunning_15000 : ");
        printNum(long_running_program(16000));
        printf("\n");
        sys_exit();
    }
    waitpid(pid6);
    sys_exit();
}
```

- **kernelMain (Changes)**
 - **Task task1(&gdt, EmptyTask);**: Creates a Task object named task1 using the GDT and assigns it the EmptyTask function. This is a igniter task. It must be added to run next tasks.
 - **Task task2(&gdt, forkExample);**: Creates a Task object named task2 using the GDT and assigns it the forkExample function.(This function explained above)
 - **taskManager.AddTask(&task1);**: Adds task1 to the task manager.
 - **taskManager.AddTask(&task2);**: Adds task2 to the task manager.

```
TaskManager taskManager;

Task task1(&gdt, EmptyTask);
Task task2(&gdt, forkExample);

taskManager.AddTask(&task1);
taskManager.AddTask(&task2);
```

System Enhancements

- **System Call Implementations**
 - New System Call Functions:
 - **sys_getpid(), sys_exec(), sys_addTask(), sys_fork(), sys_exit(), sys_waitpid().**
 - Handling System Calls:
 - The HandleInterrupt function is used to interrupt calls.
- **Multitasking Enhancements**
 - Task Management:
 - Added pIdCounter for process IDs, taskState for managing task states, and waitPid for tracking wait conditions.
 - ForkTask() duplicates the current task; AddTask() initializes new tasks.
 - Round Robin Scheduling:
 - Implemented in Schedule(), cycling through tasks based on state and timeQuantum.
 - PrintProcessTable() method prints task states.
- **Kernel Improvements**
 - Initialization and Task Management:
 - TaskManager manages tasks; forkExample initializes tasks and processes.
 - Interrupt Management

Implementation Strategies

- System Call Implementation:
 - Added system call handling methods and integrated them within the
 - InterruptManager.
- Multitasking Enhancements:
 - Improved task management and scheduling.
- Kernel Improvements:
 - Added support for multiple tasks and system calls.
 - Implemented priority-based task switching

General Flow

1. kernel.cpp uses system calls from syscalls.h
2. syscalls.cpp calls interrupts.cpp functions
3. interrupts.cpp uses multitasking.cpp functions for system calls and scheduler.